

CSCI 3150 Introduction to Operating Systems:

Assignment Two

Deadline: 18:00:00 p.m., Mon, March 6th

Total marks: 100

Topics: Semaphore and Mutex, Process Scheduling

1. 30 marks

There are M producer processes and N consumer processes sharing a buffer of size K. The interactions between them follow the rules.

- (1) Any process must access the buffer in a mutually exclusive manner;
- (2) For every piece of data put into the buffer, all consumers must receive once;
- (3) When the buffer is full, the producers must block;
- (4) The consumers must block when the buffer is empty.

Please implement a solution for such a special producer-consumer problem using wait() and post() operations with clear **c-style pseudo code** like examples in the lecture slides and explanations of the meaning of each semaphore. Your pseudo code should include **detailed comments** to ensure full understanding and avoid **potential deductions** in points.

2. 70 marks

Implement the MLFQ (Multi-Level Feedback Queue) scheduler by following the rules below (You could refer to chapter 8 of OSTEP for related information):

- **Rule 1: If $\text{Priority}(A) > \text{Priority}(B)$, A runs (B doesn't).**
- **Rule 2: If $\text{Priority}(A) = \text{Priority}(B)$, A & B run in round-robin fashion using the time slice (quantum length) of the given queue.**
- **Rule 3: When a job enters the system, it is placed at the highest priority (the topmost queue). For the jobs arriving at the same time, schedule the job with smallest pid first.**
- **Rule 4: Once a job uses up its time allotment at a given level (regardless of how many times it has given up the CPU), its priority is reduced (i.e., it moves down one queue and will be at the tail of the target queue, which means it will be scheduled last).**
- **Rule 5: After some time period S, move all the jobs in the system to the topmost queue, and sort all the jobs by pid. The job with the smallest pid will be scheduled first.**

In this question, we assume the time unit is millisecond (ms). Please note that in this assignment, we have different time allotment for different queue.

The inputs of your program: process.file and queue.cfg

- process.file contains the process information and its format is as follows:

```

ProcessNum N
pid:X1, arrival_time:X11, execution_time:X12
pid:X2, arrival_time:X21, execution_time:X22
...
pid:YN, arrival_time:YN1, execution_time:YN2

```

Here, the first line denotes there are N processes to be scheduled (there is at least one space character between ProcessNum and N), and from the second line to the (N+1)th line, each shows the pid (process id), arrival time and execution time of a process with the format like “pid:X1, arrival_time:X11, execution_time:X12” (separated by “,”). To make it simple, you can assume that all pids are different. An example is given below.

```

ProcessNum 5
pidnum:123,arrival_time:60,execution_time:90
pidnum:13,arrival_time:70,execution_time:100
pidnum:1023,arrival_time:10,execution_time:160
pidnum:12,arrival_time:80,execution_time:28
pidnum:36, arrival_time:10, execution_time:10

```

- queue.cfg contains the queue information and its format is as follows:

```

QueueNum n
Period_S S
Time_Slice_QN Xn Allotmenttime_QN Yn
....
Time_Slice_Q2 X2 Allotmenttime_QN Y2
Time_Slice_Q1 X1 Allotmenttime_QN Y1

```

Here, the first line denotes the number of queues we will use in the scheduler, the second line denotes the period S to move up all jobs to the topmost queue (Rule 5), the third line denotes the time slice for Queue n that is the topmost queue with the highest priority (the smallest time slice), and the subsequent lines define other queues with descending order in terms of priority. **Usually, the allotment time for the least priority queue will be large enough.** In each line, between the keyword and number, there should be at least one space character. One example is shown below.

```

QueueNum 3
Period_S 300
Time_Slice_Q3 10 Allotmenttime_Q3 30
Time_Slice_Q2 40 Allotmenttime_Q2 80
Time_Slice_Q1 60 Allotmenttime_Q1 120

```

The output of your program: output.log

- Your program should output the schedule to the file, output.log
- The format in output.log is as follows (This is very important as we will check the output based on this format):

Time_slot:x-y, pid:x1, arrival-time:x2, remaining_time:x3

Here, Time_slot:x-y denotes the time interval starting at time x and end at time y (Commonly the time interval is the time slice of current queue. If one process finishes without using the whole time slice, the time interval will be the time it used, and the smallest time interval is 1 ms.), pid:x1 denotes the corresponding process with pid x1 is scheduled in this time interval, arrival-time:x2 and remaining_time:x3 denote the arrival time and remaining time of the process are x2 and x3, respectively. There is at least one space character between them. You can just use the provided `outprint()` function. An example is shown below (part of the result of the given information):

```
Time_slot:10-20, pid:1023, arrival-time:10, remaining_time:150
Time_slot:20-30, pid:36, arrival-time:10, remaining_time:0
Time_slot:30-40, pid:1023, arrival-time:10, remaining_time:140
Time_slot:40-50, pid:1023, arrival-time:10, remaining_time:130
Time_slot:50-90, pid:1023, arrival-time:10, remaining_time:90
Time_slot:90-100, pid:123, arrival-time:60, remaining_time:80
Time_slot:100-110, pid:13, arrival-time:70, remaining_time:90
Time_slot:110-120, pid:12, arrival-time:80, remaining_time:18
.....
.....
```

For several situations, more specific statements are given below:

- Situation 1:

For MLFQ, if a process A finishes a time slice at time t in the top queue but not the allotment, and a new process B enters at time t. We assume that process A is enqueued first and then process B, which means process B is at the tail of this queue at time t.

- Situation 2:

For MLFQ, if a process A enters at the pre-defined Time Period, we assume to sort all processes including process A according to Rule 5.

Submission:

- **Plagiarism, incorrect file naming, and lack of comments will result in various degrees of deductions in points.**
- In question1, you need to submit the **c-style pseudo code** named '`pv_solution.c`'.
- In question2, you need to revise function **`void scheduler(Process* proc, LinkedQueue** ProcessQueue, int proc_num, int queue_num, int period)`** to implement the MLFQ (Multi-Level Feedback Queue) scheduler. You only need to submit **`scheduler-impl.c`** that can be compiled with the same "Makefile" for question2.
- TA TAN, Xin is responsible for this assignment. Questions about the assignment via Piazza are welcomed and preferred but you may also contact he via email: xtan22@cse.cuhk.edu.hk. Requests like writing code and debugging for you will be rejected according to regulations.

Attachment:

Introduction to Codes Provided in Question 2

In this document, we will present how the scheduler provided works.

1. Compile and Run

Enter your code directory and compile it as follows:

make

Then you can run it as follows:

./Scheduler

After this, you will get output like this:

```
Process number: 5
36 10 10
1023 10 160
123 60 90
13 70 100
12 80 28

Queue number: 3
Period: 300
0 60 120
1 40 80
2 10 30 /*the topmost queue*/
```

2. Implementation

Three groups of functions in the program are:

- (i) Read processes information from process.file and sort them.

```
int ReadProcessFile(); // Return the number of process in process.file, all the processes are
                        // stored in proc.
int min(int x, int y); // Return the less one between x and y.
Process MinProc(Process x, Process y); // Return the process arrive earlier; if arrive at the
                                        // same time, return the one have less pid.
void SortProcess(Process* p, int num); // Sort proc according to arrival_time and pid.
```

- (ii) Read queues information from queue.cfg file.

```
int GetQueueNum(); // Return the number of queue in queue.cfg.
int GetPeriod(); // Return the value of Period_S in queue.cfg.
void ReadQueueCfg(LinkedQueue** LQueue, int num); // Store queues into LQueue.
```

(iii) Output functions and **scheduler function that you need to implement.**

```
void InitOutputFile(); // Create void output.log file.  
void outprint(int time_x, int time_y, int pid, int arrival_time, int remaining_time);  
    // Print one line to output.log file.  
void scheduler(Process* proc, LinkedQueue** ProcessQueue, int proc_num, int
```

For above functions, if you are interested in the implementation, you can look up the source code **scheduler-exec.c**. Besides, you had better read the definition and implementation of provided data structure and functions carefully. **You can only modify the scheduler-impl.c file and add some customized helper functions in the file if needed.**

Main function:

```
int main(){
    /*
        Following codes will read and sort processes from process.file,
        and store the sorted processes into array proc[].
    */
    int proc_num = ReadProcessFile();
    Process proc[proc_num];
    for (int i = 0; i < proc_num; i++){
        proc[i].process_id = proc_tmp[i].process_id;
        proc[i].arrival_time = proc_tmp[i].arrival_time;
        proc[i].execution_time = proc_tmp[i].execution_time;
    }
    SortProcess(proc, proc_num);

    /*
        Following codes will read queues from queue.cfg, and store them into
        ProcessQueue.
    */
    int queue_num = GetQueueNum();
    int period = GetPeriod();
    LinkedQueue** ProcessQueue = (LinkedQueue**)malloc(sizeof(LinkedQueue*) *
queue_num);
    ReadQueueCfg(ProcessQueue, queue_num);

    /*
        Initiate output.log file.
    */
    InitOutputFile();

    //Call scheduler() here.
    scheduler(proc, ProcessQueue, proc_num, queue_num, period);

    return 0;
}
```